

The Noisy Work of Uncertainty Visualisation

Harriet Mason¹ Dianne Cook¹ Sarah Goodwin²
Emi Tanaka³ Susan VanderPlas⁴

¹Department of Econometrics and Business Statistics, Monash University,
Melbourne, Australia,

²Department of Human Centred Computing, Monash University, Melbourne,
Australia,

³Biological Data Science Institute, The Australian National University, Canberra,
Australia,

⁴Department of Mathematics & Statistics, Utah State University, Logan, UT,
United States,

Annual Review of Statistics and Its
Application 2023. AA:1–25

[https://doi.org/10.1146/annurev-
statistics-XXXXXX-XXXXXX](https://doi.org/10.1146/annurev-statistics-XXXXXX-XXXXXX)

Copyright © 2023 by the author(s).
All rights reserved

Keywords

data visualisation, statistical graphics, data science

Abstract

Better representation of the uncertainty in a data visualisation is a focus of recent research activity. A problem with the current literature is that there is a lack of clarity about the definition of uncertainty and what it means to represent it in a plot. This confusion has led to conflicting results in the literature, especially in experiments that assess the effectiveness of different uncertainty representations. In this review, we summarise the current literature, provide workable definitions, and illustrate these definitions with examples. In doing so, we ask what it really takes to achieve transparency in statistical graphics. It is hoped that it will be useful for guiding new graphics methodology and experimental research.

1. Introduction

What do we mean when we talk about “uncertainty visualisation”? The phrase can feel contradictory to anyone familiar with the individual terms. Among statisticians, “uncertainty” is often discussed as an omnipresent spectre touching every stage of our analysis without ever being fully seen. Authors will often mention that the phrase is vague ([Spiegelhalter 2017](#), [Griethe & Schumann 2006](#)), or avoid defining it by describing a list of things uncertainty *could* be ([Kinkeldey et al. 2014](#), [Hullman 2016](#)), but rarely do authors attempt to discuss what uncertainty actually *is*. By contrast, visual statistics (which includes information visualisations and data plots) are one of the most powerful tools in the statistician’s toolbox, allowing for quick and memorable communication that identifies quirks in our data that we didn’t even know to look for. We see this in datasets such as Anscombe’s quartet ([Anscombe 1973](#)) or the Datasaurus Dozen ([Matejka & Fitzmaurice 2017](#), [Locke & D’Agostino McGowan 2018](#)), where visual statistics are able to highlight elements of the data that are invisible to the typical summary statistics. We also see this in recall experiments, where simply sketching a distribution before recalling statistics or making predictions can greatly increase the accuracy of those measures ([Hullman et al. 2018](#), [Goldstein & Rothschild 2014](#)). Visualisation of uncertainty enables us to pull back the curtain and explore the unknowns of our analysis.

As nice as this sentiment is, it turns out to be easier said than done. Reviews on uncertainty visualisation rarely offer tried and tested rules, and tend to focus on the difficulties faced when trying to summarise the field. [Kinkeldey et al. \(2014\)](#) found most experimental methods to be ad hoc, with no commonly agreed upon methodology, formalisation, or description of general principles. [Hullman \(2016\)](#) noticed there is a serious noise issue in the field, with errors from participants misunderstanding visualisations, misinterpreting questions, and incorrectly applying heuristics, overwhelming any information we can glean from studies. [MacEachren et al. \(2005\)](#) identified so many contradictions that they spent an entire page discussing the conflicting evidence for the question “Should I map uncertainty to colour hue?”. [Spiegelhalter \(2017\)](#) concluded that different plots are good for different things, arguing against a universal best plot for all people and circumstances. [Padilla et al. \(2022b\)](#) summarised several cognitive effects that repeatedly arise in uncertainty visualisation experiments, but these effects were each discussed in isolation as a list of considerations rather than an overarching theory for effective uncertainty visualisation.

“Science is built up of facts, as a house is built of stones; but an accumulation of facts is no more a science than a heap of stones is a house.” - Henri Poincaré (1905)

While these reviews are thorough in scope, none discuss how the existing literature contributes to the broader goal of uncertainty visualisation – that is, despite the wealth of reviews, the field of uncertainty visualisation remains a heap of stones. There is a mountain of work that identifies common heuristics found in uncertainty visualisations, evaluates competing plot designs, or starts a theoretical discussion on a niche aspect of the field. While important, each of these papers offers up its own bespoke motivation and methodology, with little reference to the uncertainty visualisation papers outside their fiefdom. The field is in desperate need of a unifying theory that can tie the conflicting and siloed research together. This review attempts to address this issue by offering a novel perspective on the uncertainty visualisation problem. That is, we will use the wealth of established stone to construct a foundation to build a house.

2. The purpose of uncertainty visualisation

Mentions of “uncertainty visualisation” start springing up around 1990, across several different fields (Ibrekk & Morgan 1987, MacEachren 1992), each with its own motivation for the work. In computer science, the area appears to be motivated by issues in the public’s perception of random variables, with the hope that visualisations would give laypeople the ability to extract important information from graphical representations (Ibrekk & Morgan 1987). With similar concerns about the public’s understanding of randomness, the fields of psychology, statistics, and economics used “uncertainty visualisations” as a communication tool to mitigate the psychological bias associated with the communication of risk, a topic of concern since the early 1980s (Spiegelhalter 2017). In cartography, it was motivated by the inherent uncertainty of geoscience data, the practical use of visualisation as an exploratory tool, and the constrained visual channels from map representations (MacEachren 1992). These disparate motivations have blended together, and today, uncertainty visualisation is usually motivated by the vague goal of “decision-making”. This term has been used to mean the mitigation of psychological bias to ensure economically rational decisions (Padilla et al. 2022a, 2021, Kale et al. 2021), the facilitation of trust or confidence (Zhao et al. 2023, Yang et al. 2023), the ability to extract values related to a distribution (Sarma et al. 2023), the prevention of false discovery in plots (Sarma et al. 2024, Koonchanok et al. 2023), or the extraction of some other metric that is vaguely related to “uncertainty” (Chakraborty et al. 2024, Ndlovu et al. 2023).

Given that there is so much disagreement in uncertainty visualisation, how did these disparate motivations come to be seen as interchangeable? All discussions on uncertainty visualisation seem to have a common thread that connects them: the belief that the ultimate goal of uncertainty visualisation is not trust, rationality, or value extraction, but *transparency*. We see it said directly in reviews of the field (Padilla et al. 2022b), or when authors claim that failing to include uncertainty is akin to fraud or lying (Hullman 2020, Manski 2020). We see it when authors assert that uncertainty communicates the legitimacy (or illegitimacy) of the conclusion drawn from visual inference (Kale et al. 2018, Griethe & Schumann 2006). We see it when authors say uncertainty visualisations should communicate a degree of confidence (Correll et al. 2018, Boukhelifa et al. 2012) or validity (Hullman 2020, Griethe & Schumann 2006) in our conclusions. We see it when authors suggest uncertainty visualisation should “guide, qualify, or soften our judgements of uncertain data” (e.g., Wilkinson (2005) in his seminal work on the grammar of graphics). These authors are not wrong about the need for transparency in science communication: a six-month survey of anti-mask groups on Facebook during the COVID-19 pandemic showed that anti-maskers made persuasive arguments by exploiting inherent uncertainty ignored by pro-maskers (Lee et al. 2021). Therefore, if transparency is necessary and important, what does it look like?

Uncertainty visualisation is motivated by the need for a sort of “visual hypothesis test”, a sentiment expressed by some authors directly (Correll & Gleicher 2014, MacEachren 1992). A successful uncertainty visualisation would act as a “statistical hedge” for any inference we make using the graphic. Since the purpose of a visualisation is to give a quick gist of the information (Spiegelhalter 2017), this hedging should be communicated visually without the need for complicated mental calculations. Therefore, an effective uncertainty visualisation should not just “show” uncertainty; untrustworthy conclusions should *not be visible*. If we refer to the conclusion we draw from a graphic as its “signal”, and the uncertainty that should make this signal harder to read as the “noise”, we can summarise the above

information into three key requirements. A good uncertainty visualisation should perform:

- 1) Signal enhancement: reinforcing statistically significant signals to encourage confidence in results, and
- 2) Signal suppression: hiding spurious signals that are overwhelmed by noise.

Additionally, the performance of tasks 1) and 2) should be proportional to the level of confidence in those conclusions. When correctly performing uncertainty visualisation, we should feel reassured that what we see is actually there. Therefore, we **suggest calling this approach to uncertainty visualisation “signal modulation”**. This is the main criterion we will use to assess the current literature on uncertainty visualisation.

3. Current Approaches

3.1. Ignoring uncertainty

The most common way to visualise uncertainty is to simply not. A study conducted by [Hullman \(2020\)](#) found that only a quarter of authors surveyed included uncertainty in 50% or more of their visualisations, in part because authors are not sure how to calculate uncertainty. This is not entirely unreasonable, given that even uncertainty visualisation researchers themselves seem to be in conflict about what exactly uncertainty is. We will start with visualisations that ignore uncertainty with the hope that by looking at where uncertainty isn’t, we can better understand where it is.

3.1.1. What is uncertainty? It is surprisingly hard to describe what uncertainty is. Most authors avoid the problem and describe the many characteristics of uncertainty. Often, uncertainty is split by factors such as whether it is due to true randomness or a lack of knowledge ([Begg et al. 2014](#), [Spiegelhalter 2017](#), [Gustafson & Rice 2019](#), [Padilla et al. 2022b](#), [Hullman 2016](#), [Walker et al. 2003](#)); quantifiable or unquantifiable ([Spiegelhalter 2017](#), [Walker et al. 2003](#), [Padilla et al. 2022b](#)); scientific or human ([Benjamin & Budescu 2018](#), [Gustafson & Rice 2019](#)); systematic or random ([Sanyal et al. 2009](#)); statistical or bounded ([Gschwandtner et al. 2016](#), [Olston & Mackinlay 2002](#)); due to a lack of accuracy or precision ([Griethe & Schumann 2006](#), [Benjamin & Budescu 2018](#), [Hullman 2016](#)); etc. There are enough qualitative descriptors of uncertainty to fill a paper, but none of this is particularly helpful in understanding how to integrate it into a visualisation.

Rather than trying to define uncertainty by looking at the myriad ways in which it *does* appear in an analysis, we may find it easier to look at where it *does not*. Descriptive statistics describe our sample as it is and summarise large data into a usable format, but they are not seen as the primary goal of modern statistics. In 19th-century England, *positivism* was the popular philosophical approach to science. Practitioners of the approach believed statistics ended with descriptive statistics, as science must be based on actual experience and observations ([Otsuka 2023](#)). Positivists included famous statisticians such as Francis Galton and Karl Pearson. In order to make statements about population statistics, future values, or new observations, we need to perform inference, which requires the assumption of the “uniformity of nature”, that is, we need to assume that unobserved phenomena should be similar to observed phenomena ([Otsuka 2023](#)). Positivists believed that referencing the unobservable was bad science, and that one should not make claims beyond descriptive statistics. Since uncertainty is non-existent in descriptive statistics, it is clear that uncertainty

is a by-product of inference: uncertainty is the noise that is both inseparable from our inference and meaningless without it. If we consider uncertainty to be a by-product of statistical inference, then uncertainty visualisations are the plots that depict an estimate, and therefore have an associated uncertainty. The most complete description of these estimates is their distributions.

Rather than extracting just one element of the distribution, if you can retain the whole distribution, that not only allows the uncertainty calculation to be reproduced, but also makes it possible to derive other estimates as well. Suggesting distributions as a representation of uncertainty is not new. [Kay \(2023\)](#) originally suggested thinking about uncertainty visualisations as visualisations with distribution inputs, to replace the commonly used mean and standard deviation and remove the assumption of a Gaussian distribution. In practice, uncertainties are sometimes calculated by other people or organisations, and the process used to derive them may not be known. While some researchers believe these abstract notions of uncertainty, such as credibility ([Thomson et al. 2005](#)), forecaster confidence ([Padilla et al. 2021](#)), or uncertainty about uncertainty ([Hadjimichael et al. 2024](#)), are too complex to be quantified, this is not necessarily true. Abstract concepts such as human belief or credibility are regularly quantified by Bayesians, and hierarchical approaches are often used to model uncertainty about uncertainty.

3.1.2. Example: ignoring uncertainty. If visualising uncertainty is fundamentally visualising a set of random variables, what does “ignoring” uncertainty look like? Figure 1 shows plots of data from three different scenarios, each with either high or low uncertainty. The plots show the expected value of the input distribution. In plots R1 and R2, the expected value is the point estimate from a simple linear regression on a car’s miles per gallon (`mpg`) and weight (`wt`) using the `mtcars` data (available in `ggplot2` ([Wickham 2010](#)), and originally from [Henderson & Velleman \(1981\)](#)). For the low uncertainty case, the linear regression is calculated on the full `mtcars` data, but in the high uncertainty case, it is calculated using a subset of three points. The points used to calculate the linear regression are also shown on the plot; however, the fitted line is the only component that should include any uncertainty. Plots S1 and S2 show a choropleth map of Iowa, where counties are coloured according to a simulated temperature measurement, with measurement error being the uncertainty associated with the measuring instrument. In S1 the measurement error is low, while in S2 the measurement error is high. Plots G1 and G2 show samples simulated from five populations (A-E). In G1, the distributions have a low variance, and in G2, they have a high variance. In the linear regression, we see a downward trend, in the map, we see a sine wave, and in the univariate distributions, an incremental increase in treatment effect. The noise in the high uncertainty case is set such that it should overwhelm the trend in all three plots. Can you see a difference between the plots in the top row versus the bottom row? Is the strength of the trend communicated through the visualisation? The answer to both of these questions, for S1, S2, G1, and G2, is no as the high and low variance cases are identical. For R1 and R2, any ability to differentiate the plots would likely relate to the data points, rather than the regression line. The lack of differentiability between the high and low uncertainty cases highlights the danger of ignoring uncertainty in our visualisations.

These examples will serve as touchstones for a discussion of uncertainty visualisation, focusing on the approaches suggested in the literature.




Figure 1: Three example types of data and associated plots that will be used to illustrate various choices of uncertainty representation throughout the paper: scatterplot and regression line (R1, R2), spatial choropleth map (S1, S2), and grouped dotplot (G1, G2). Two levels of uncertainty (low, high) are used with each example. Ignoring the points in R1 and R2, there are no differences between the high and low uncertainty versions. Ignoring uncertainty can lead to misrepresentation of data.

3.2. Uncertainty as a statistic

To avoid ignoring uncertainty entirely, some authors opt to visualise uncertainty using an “uncertainty statistic”. Rather than visualise the statistic we wish to draw inference on, we calculate some statistic related to the uncertainty (such as a probability, variance, or error) and visualise that instead. For example, in geospatial visualisation, we might see an exceedance probability map (Kuhnert et al. 2018), where areas are coloured according to the probability that the estimate exceeds a specified value. In risk communication, we might see an icon array (or pictograph) (Spiegelhalter 2017) that communicates the relative number of affected individuals for a particular treatment using coloured human icons. In statistics, we have the summary plot of Potter et al. (2010), which combines the plot of a single set of numeric values with a density estimate, boxplot, and statistical moments.

Authors take these approaches because they explicitly see uncertainty as a variable of importance in and of itself (Blenkinsop et al. 2000). This is a perfectly acceptable motivation as uncertainty is often a variable of importance, but should we consider these visualisations “uncertainty visualisations”? Is making an uncertainty visualisation that easy?

3.2.1. Example: visualising variance. Figure 2 depicts the six plots introduced in Figure 1, but instead of visualising the central value of the distribution, we visualise an uncertainty statistic. Plots R1 and R2 show the residual plot of our linear regression, which is a common

approach for isolating and evaluating the noise in a regression. By changing the underlying statistic, we have completely changed the interpretation of our plot, and therefore, the plot cannot be evaluated using our signal modulation criteria. While the divergence from the signal modulation goal is clear in the case of regression residuals, this is not always the case. Plots S1 and S2 show an exceed probability map, where each county is coloured according to $\hat{P}(\text{temperature} > 27)$. At first glance, we might think this plot is performing signal modulation as the sine wave trend is clearly visible when the error is low, but barely visible with high error; however, just like R1 and R2, the meaning of these plots was changed when we swapped the underlying statistic. As a temperature of 27 is a middling value in our example, as the variance increases, $\hat{P}(\text{temperature} > 27)$ approaches 0.5 for all counties. Exceedance probability maps are often used to identify extreme values, and if we replaced 27 with a more extreme value (such as 35), the signal would swap, and we would see the wave pattern in the high uncertainty case, but not the low uncertainty case. Continuing this trend, G1 and G2 visualise the variance of the distributions, and once again show that changing the underlying statistic changes the fundamental meaning of our graphics, but with a slight twist. This plot reveals some interesting information about our data: the variance is constant in G1 but different for each group in G2. Here, we highlight what is *actually* going on when we visualise uncertainty as a statistic. The approach allows us to glean new information about the variance itself, but this information is largely independent of the signal conveyed by Figure 1.

3.2.2. What is an uncertainty visualisation, then? What an uncertainty visualisation is or is not is one of the most pervasive divides in the literature. For example, [Wilkinson \(2005\)](#) mentions that popular graphics, such as pie charts and bar charts, omit uncertainty. [Wickham & Hofmann \(2011\)](#) suggests their product plot framework, for area plots like bar charts and also histograms, needs to be extended to include uncertainty representation. However, pie charts, bar charts, and histograms have all been used in a significant number of experiments as examples of an “uncertainty visualisation” ([Ibrekk & Morgan 1987](#), [Olston & Mackinlay 2002](#), [Zhao et al. 2023](#), [Hofmann et al. 2012](#)). These examples imply that a chart being classified as an “uncertainty visualisation” is entirely at the discretion of the author, which brings into question whether the category is meaningful at all. How did such a conflict arise in the literature?

This conflict stems from an unrecognised disagreement about the purpose of uncertainty visualisation. If you think uncertainty visualisation is about communicating risks or probabilities, uncertainty visualisations are just visualisations of “uncertainty statistics”. On the other hand, if you think uncertainty visualisation is about signal modulation, then you see an uncertainty visualisation as a transformation of an existing graphic that adds the uncertainty in. The former has no limitation on the visual appearance of an “uncertainty visualisation”, allowing pie charts, bar charts, or histograms, so long as the graphic is visualising “uncertainty”, while the latter believes uncertainty visualisations only exist in relation to some “normal” visualisation.

While arguments have been made for both approaches, there are two key reasons we shouldn’t think of visualisations of “uncertainty statistics” as uncertainty visualisations. The first problem is that it goes against established norms in the field of visualisation. When we refer to the graphics depicted in Figure 2 as “uncertainty visualisations”, we are classifying visualisations by the data they display, not by their visual features. A scatter plot that




Figure 2: Treating the uncertainty as a statistic, with the same six examples. The regression (R1 and R2) is a scatterplot of residuals vs explanatory variable, separated from the regression, making comparison of uncertainty related to the trend more difficult. For the choropleth (S1, S2): instead of temperature, the probability of exceeding 27°C is shown. This has the effect of highlighting (sine wave) trend in the low error data, and de-emphasising it in the high error data. The plots G1 and G2 have replaced the treatment effect with the variance of our treatment effect. It is a bit nonsensical, but we learn something interesting that was not seen earlier: the variance in G2 is not constant like it is in G1.

compares means and a scatter plot that compares variances are both scatter plots. There is no reason for “uncertainty statistics” to deviate from these conventions. The second problem is slightly more subtle, and that is that the “uncertainty” in these graphics is not acting as uncertainty at all; it is acting as signal.

How do we know that the uncertainty in these visualisations is acting as signal and not as noise? The role of signal and noise is not set in stone and can change depending on the goals of our analysis. [Meng \(2014\)](#) commented that what is kept as data and what is tossed away is determined by the motivation of an analysis - what was previously noise can become signal depending on the question. [Otsuka \(2023\)](#) argued that the entire process of boiling real-world entities down into data requires us to have an internal model, which he refers to as “probabilistic kinds”, that are defined by the relationships we seek to identify. [Wallsten et al. \(1997\)](#) argue that the best method for evaluating or combining subjective probabilities depends on the uncertainty the decision-maker wants to represent, and why it matters. [Fischhoff & Davis \(2014\)](#) suggested we should have methods for communicating uncertainty depending on what the user is supposed to do with it. [Spiegelhalter \(2017\)](#) says we “cannot assess the quality of risk communication unless the objectives are clear”. [Peña-Araya et al. \(2025\)](#) asserted that whether or not uncertainty is a source of doubt depends on the context.

The sentiment behind these discussions is clear: the role of signal and noise is determined by the hidden parameter we are trying to identify; whether that parameter is a measure of central tendency or spread is largely irrelevant. When we swap a linear regression out for a plot of our residuals, the goal of the visualisation, and therefore the roles of signal and noise, change as well.

If the uncertainty in a graphic is acting as a signal, there isn't an interesting perceptual challenge associated with the visualisation: the uncertainty can be displayed using standard principles of graphic design. In changing the inferential statistic, we also haven't dealt with the original problem of integrating noise, as these "uncertainty statistics" *also have associated uncertainty in the estimates* (e.g., variance of standard deviation estimate) that is being ignored. There is nothing wrong with explicitly visualising variance, error, bias, or any other statistic. These metrics provide important and useful information for analysis and decisions. The problem with this approach is that it means everything is an uncertainty visualisation, and if everything is an uncertainty visualisation, nothing is.

3.3. Uncertainty as a variable

It is reasonable to assume that the "uncertainty as a statistic" approach falls short because we did not visualise the uncertainty within the context of our estimate. This hypothesis leads to another common approach to uncertainty visualisation, where uncertainty is seen as an additional variable that needs to be integrated into the original visualisation of our estimates. Under this philosophy, uncertainty visualisation becomes a question of how we can add additional information (uncertainty) into an existing visualisation without cluttering the plot and overwhelming the viewer (Griethe & Schumann 2006, Peña-Araya et al. 2025). This approach emerges in computer science (Kinkeldey et al. 2014), cartography (MacEachren et al. 2005), and statistical graphics (Wilkinson 2005) alike. This "free aesthetic" approach to uncertainty visualisation encourages us to map any standard error to any unused aesthetic, with the primary concern being the intuitiveness of the mapping (Maceachren et al. 2012).

When uncertainty is regarded as a variable, discussions focus on how integrated the uncertainty is with the estimate. Kinkeldey et al. (2014) identified a split between intrinsic plots, where we map uncertainty to colour or size of the geometric object of our estimate, and extrinsic plots, where uncertainty is mapped to a separate geometric object, such as glyphs or error bars. Similarly, Padilla et al. (2022b) classified uncertainty visualisations as graphical annotations (extrinsic), probability mapped to a visual encoding channel (intrinsic), or a hybrid of the two. We can reframe these two approaches within the grammar of graphics. In the grammar, *geometries* are the geometric objects we draw on the plot, while *aesthetics* represent things like position, colour, and size of said geometric objects. In this framework, extrinsic approaches manage uncertainty at the geometry stage, while intrinsic approaches manage uncertainty at the aesthetic stage.

While the intrinsic versus extrinsic dichotomy is interesting, it is unclear if these levels of "integration" have any effect on the plot's ability to perform signal modulation.

3.3.1. Example: mapping two variables. Figure 3 shows the six examples introduced in Figure 1, with uncertainty mapped using an intrinsic or extrinsic approach.

In plots R1 and R2, we take an intrinsic approach to visualising uncertainty, where the

estimated standard error of the slope is represented by the width of the line, while the error of the intercept is mapped to transparency. While these help to give a gestalt of the uncertainty, they are not exacting representations. Typically, we would visualise the uncertainty in a linear regression using a ribbon geometry, which would allow the uncertainty to change with respect to the wt variable. By opting to use an intrinsic method, we highlight a key limitation of this approach: it assumes the uncertainty is constant for each geometric object. Not only does this prevent the uncertainty from varying with x , but it also prevents us from displaying multiple standard errors. Here we have used one standard error, but why not map the width to two standard errors, or even three? The obfuscation of the variance in the slope is not the only issue with this intrinsic mapping. We mapped the standard error of the intercept to transparency, as it is a highly intuitive mapping; however, translating this transparency to a numerical quantity is virtually impossible. While this approach renders this aesthetic mapping effectively useless, this is the mapping that this design approach suggests.

Another intrinsic approach is illustrated in the bivariate colour palette maps shown in plots S1 and S2. With two dimensions of the plot reserved for spatial position, it is difficult to incorporate the error. The bivariate colour palette maps the estimate to hue and the error to the saturation and brightness, keeping the signal and noise contained to one visual aesthetic. In this case, even though the signal and noise are mapped within a single visual channel (colour), the wave pattern remains visible in both plots.

The extrinsic approach, shown in plots G1 and G2, has the estimate represented by a point, and uncertainty computed as a 95% confidence interval mapped to line length. Our plot now clearly shows the change in variance in the high uncertainty case, but the incremental trend in the average is still visible. Therefore, the display fails to sufficiently suppress the high uncertainty signal.

We can see that all three approaches are unsuccessful when we evaluate their ability to perform signal modulation. Because all of these graphics visualise the distribution's estimate and uncertainty as two separate pieces of information, the message of the plot is "here is the trend *and* here is the uncertainty". It is worthwhile to examine why this occurs, to see if we can move towards a version of this plot where we are able to communicate signal and noise simultaneously.

3.3.2. Can we visualise a "single integrated uncertain value"? The reality is, based on our discussion on inferential statistics, uncertainty *isn't* a separate variable: it is a component of the random variable that is indistinguishable from the random variable itself. Similarities between the "as a variable" approach and the "as a statistic" approach become apparent when investigating why some authors opt for the variable approach in the first place. When visualising a point estimate (i.e., Figure 1) and variance (i.e., Figure 2) side-by-side, one needs to switch focus between the two displays, which makes it difficult to combine the information on the graphics (Moritz et al. 2017, Correll et al. 2018). This implies that the reason authors stop visualising uncertainty as a statistic and instead visualise uncertainty as a variable is because of the cognitive burden imposed by switching focus, rather than the understanding that the method is not philosophically sound. To achieve signal modulation, we need to visualise noise and signal together as a "single integrated uncertain value" (Kinkeldey et al. 2014) rather than as two separate statistics.




Figure 3: Treating the uncertainty as a variable, with the same six examples. In plots R1 and R2, the standard error of the slope is mapped to the line width. The standard error of the intercept is mapped to the transparency, which is less conventional. In plots S1 and S2, a bivariate colour palette is used with mean mapped to the hue, and estimated standard error mapped to saturation and brightness. In plots G1 and G2, the mean of each group is visualised as a point, and the standard error is visualised as an interval. In this example, the uncertainty has been integrated with the signal, but the results are undesirable. This is especially true in the case of the choropleth map, where the signal is easier to see when the error is higher.

Just because we can still see the signal in Figure 3, that does not mean the reading of the estimate is completely independent of the uncertainty. When making any visualisation, we usually want the visual channels to be separable, that is, we don't want the data represented through one visual channel to interfere with the others (Smart & Szafr 2019). Separability may be desirable in standard data visualisation, but in uncertainty visualisation, it allows the estimate and its variance to be read independently, potentially leading to the uncertainty being ignored (Padilla et al. 2022b). Therefore, rather than trying to maintain visual separability, the goals of uncertainty visualisation align far better with the pursuit of visual integration. In an ideal system, our estimate and uncertainty would be manipulated separately, but would be so *well-integrated that they are read as a single channel by the human brain*. The problem is that even when mapping uncertainty and estimate to the same aesthetic, as illustrated by the bivariate colour palette map in Figure 3, our methods fall short. Colour hue and brightness are classic examples of integrable variables (Vanderplas et al. 2020), and decreasing saturation should make the colours harder to distinguish, but the signal is still clearly visible when the variance is large. Worse, multi-dimensional colour palettes can make the graphics harder to read and less accessible (VanderPlas & Hofmann

2015).

3.3.3. Mapping combined variables. The Value Suppressing Uncertainty Palette (VSUP) (Correll et al. 2018) was designed with the intention of preventing high uncertainty values from being extracted from a map, and this is achieved by blending colours together as they become less certain. Figure 4 shows the spatial example (S1, S2) using this approach. Since the palette was designed with the extraction of individual values in mind, it has only been tested on simple value extraction tasks (Correll et al. 2018) or search tasks (Ndlovu et al. 2023). We can see that, at least for our example, when the uncertainty is high, the spatial trend has functionally disappeared.




Figure 4: The spatial examples displayed with a choropleth map using a VSUP colour palette, where hue is blended when uncertainty is increased. Plot B has successfully produced signal suppression. Look closely at the scales, though: it may have another explanation.

Finally, we have signal suppression! Well, not really, sorry, we tricked you. The two plots depicted in Figure 4 actually show the exact same data; they are both the low variance case. If you look closely, you can see that the two plots have a different scale. Plot A has been scaled manually using the maximum and minimum error from the high and low uncertainty cases, while plot B has been scaled automatically using only the data from the low uncertainty case, which is the only data that would be visible in the plot. Since the variance and estimate are scaled independently, arbitrary differences in the range of our variance, unrelated to the estimate itself, will have significant impacts on the visual appearance of our plot. The scale issue in VSUP maps was also recognised by Kay (2019), who noted that the suppression of any one hypothesis largely depends on the methods we use to combine the palette, and the variance levels at which the blending occurs. This means that, for us to know that our plot will successfully perform signal validation, we need to already know what signal we are trying to validate and set up the VSUP palette accordingly. This means that VSUP maps are not suitable for exploratory data analysis.

3.3.4. Uncertainty and exploratory data analysis. The lack of uncertainty in descriptive statistics is due to the lack of inference. Descriptive statistics are actually a small piece of a much larger field, exploratory data analysis (EDA), that tends to de-emphasise statistical inference. Tukey et al. (1977) described EDA as the process of searching for interesting hypotheses (“the greatest value of a picture is when it forces us to notice what we never expected to see”), and defined it in relation to confirmatory data analysis (CDA), the process of verifying a hypothesis. Subfields of EDA include: initial data analysis (Huebner

et al. 2016, Chatfield 1985), which involves checking assumptions and data quality prior to CDA, and model diagnostics (e.g., Belsley et al. (1980)), including posterior checks of model fit. What binds these pursuits together is their reliance on visual summaries for making assessments and an absence of formal inference.

Hullman & Gelman (2021) argued that the EDA and CDA are not entirely distinct, as it is often difficult to draw a hard line. Our belief is that these exist on a continuum, where we have an inherent trade-off between the number of hypotheses we can look for and the certainty of any conclusions reached. It can help to think of the knowledge-generating process of EDA and CDA as the nozzle on a hose with multiple spray options, where EDA is a fine misting spray that touches everything in the room, and CDA is a high-pressure jet capable of obliterating any and all debris from any single spot.

Viewing EDA and CDA as a dichotomy can create some confusion when it comes to understanding the source of uncertainty in our analysis. This is why we have avoided the topic until now, despite the fact that an uncertainty visualisation system for EDA is one of the most discussed topics in the field (Hadjimichael et al. 2024, Peña-Araya et al. 2025, MacEachren et al. 2005, Sarma et al. 2024, Griethe & Schumann 2006). The EDA versus CDA dichotomy can be compared to the dichotomy between induction, for building theories, and deduction, for testing theories, from formal logic. One of the trade-offs in the two methods is that deductive conclusions provide certainty, while conclusions from induction are inherently uncertain. This means that EDA, the inductive counterpart, has uncertainty in any conclusions reached, with the requirement to follow up our newfound hypothesis with CDA if we want true certainty. This adds another layer of confusion to the study of uncertainty visualisation: conflating the uncertainty in our data with the uncertainty that is inherent to an exploratory process (EDA).

Authors often over-compensate for the inherent EDA uncertainty, pre-emptively hedging against every false inference that could possibly be drawn from a graphic. Hullman & Gelman (2021) argues there is no such thing as a “model-free” visualisation. Guo et al. (2025) provides a grammar for visualising statistical model checks. The lineup protocol (Buja et al. 2009) provides the viewer with plots of the data in a field of plots of null data where any patterns seen are due to sampling variability. The Rorschach protocol, from the same paper, shows only null plots to give the reader some intuition for what spurious sampling patterns exist. Savvides et al. (2019) provides a statistical super-test against multiple comparisons driven by probabilistic arguments. There is a CDA quality to these approaches. The sum total of a lot of CDA is not EDA, just as swinging a high-pressure jet around a room is not equivalent to using a misting spray. While EDA and CDA may be along a continuum, this comes from the large number of hypotheses that can be identified using visualisation. For any particular hypothesis, we cannot simultaneously perform EDA and CDA, as doing so would require us to violate causality. It is impossible to validate a hypothesis before we even know it exists.

To truly create an uncertainty visualisation approach that is capable of EDA, we need to accept that uncertainty is inherent to the method, and it cannot be pre-emptively removed from the visualisations. If we accept this fact, then the only possible source of uncertainty in an uncertainty visualisation system that performs EDA is from the data itself. That is to say, the only possible way for there to be uncertainty in a visualisation designed for EDA is that the data itself is a summary resulting from inference that has been done earlier in our

analysis. We can see this in the original description of Figure 1, where the uncertainty in all cases, even measurement error, represents inference that was performed earlier in our analysis.

In the VSUP approach, our ability to arbitrarily decide which values to blend at, or which suppression approach to use, means that the “uncertainty” we are visualising will be informed by the conclusions we are drawing, and not a product of the data itself. This is an antithetical approach to EDA. For a visualisation to be suitable for EDA, it should always look the same regardless of what hypothesis we plan to draw from it. As much of the transparency in data visualisation comes from this feature in EDA, it is reasonable to set it as a requirement of our visualisations. This means we cannot treat signal and noise as separate variables, but rather we must treat them as a single integrated unit.

3.4. Uncertainty as a distribution

Rather than trying to reduce random variables to a single value or pair of values, why not visualise the whole distribution? This approach is found in hypothetical outcome plots (HOPs), which animate a sequence of potential outcomes of a distribution (Hullman et al. 2015), in pixel-maps (Lucchesi et al. 2021, Blenkinsop et al. 2000), and in visualisations of multiple forecasts (Hyndman & Athanasopoulos 2021).

3.4.1. Example: visualising samples. Figure 5 shows the data from the six examples introduced in Figure 1, as distributions. The distribution is represented by either a sample of outcomes or a quantile dot plot. Plots R1 and R2 show a linear regression as a sample of possible outcomes from the estimated bivariate normal distributions of the slope and intercept. Plots S1 and S2 show a pixel map, which is a choropleth map where each county is subdivided into an $n \times n$ grid. Each grid cell is coloured according to a simulated outcome from the county’s distribution, giving the map a pixelated appearance. Plots G1 and G2 show each univariate distribution as a quantile dot plot. Quantile dotplots plot a uniform sample of quantiles as a dotplot, offering a discrete alternative to density plots (Kay 2023). We can see that the strong downward trend in the linear regression, the sine wave in the choropleth map, and the incremental increase in the univariate distributions are all clearly visible in the low variance case, but disappear in the high variance cases. The graphics have achieved signal validation. Visualising the estimate as a distribution also gives additional information, such as the previously hidden end peaks of the univariate distributions in G1 and G2.

3.4.2. Quantified versus unquantified uncertainty. By showing our distribution as “data”, we are able to read the “uncertainty” plots (i.e., Figure 5) using the same perceptual mechanisms we use to read the “non-uncertainty” plots (i.e., Figure 1). This should lead to more effective communication, as people tend to read more complicated visualisations like the bivariate and VSUP plots the same way they read the simple choropleth counterpart (Ndlovu et al. 2023). This approach also does not significantly hinder our ability to extract the individual values mapped by the previous plots, as extracting global statistics from a sample can be done with relative ease (Franconeri 2021). For example, in Figure 5, even though we have not included the average in G1 or G2, a human viewer can still extract the average effect of each group with relative ease. We are able to include more information by offloading more computation to visual processing, but what is the limitation of this approach? Which aspects




Figure 5: Visualising the uncertainty as draws from an estimated distribution, with the same six examples. Plots R1 and R2 show a linear regression as a sample of possible outcomes from the distribution. Plots S1 and S2 show a pixel map. Plots G1 and G2 show the groups as quantile dot plots. In each case, the signal (regression line, sine wave, increasing trend) has disappeared with high uncertainty.

of uncertainty should be processed by our visual system, and which should be processed by our statistical computation? This is actually a question about how much of our uncertainty should be quantified.

Quantified uncertainty usually focuses narrowly on concepts that can be calculated from theoretical or empirical distributions, such as probability, confidence intervals, variance, error or precision (Hullman et al. 2018, Maceachren et al. 2012, Thomson et al. 2005, Kay 2023). On the other hand, unquantified uncertainty often includes a broader range of concepts like missing values, reliability, model validity, or source integrity (Griethe & Schumann 2006, Wilkinson 2005, Pang et al. 1997, Pham et al. 2009, Boukhelifa et al. 2017). When discussing uncertainty, we typically consider these unquantified uncertainties, not because these things *are* uncertainty, but because they can *create* uncertainty when we perform inference. This is often because these unquantified uncertainties violate our model assumptions. Quantifying the effect of these assumption violations on our estimates transforms unquantified uncertainty into quantified uncertainty.

Sometimes we are able to visualise these assumption violations directly. For example, we can check for structure in our missing data using the `naniar` package Tierney & Cook (2023) that allows us to include missing values as a “shadow” alongside our usual visualisations. This approach amounts to just “showing the data”, which is a simple but effective option for uncertainty visualisation that is largely overlooked. While this approach is useful for better understanding data, it will not eliminate visual signals that are rendered invalid due to

incorrect assumptions about the model or missing data. We can only integrate uncertainty as noise when that uncertainty has been *quantified* as an effect on the estimates we are visualising. This is not to say one method is preferable; visualising both quantified and unquantified uncertainty is necessary for a healthy analysis. Data analysis often works in cycles, where we find assumption violations using EDA, quantify the effect of these violations on inference, and then visualise the output of that inference using uncertainty visualisation.

4. Evaluating uncertainty visualisations

Unfortunately, the conflicting results in the field are not limited to plot design and extend to the experimental findings as well. Reviews that attempt to synthesise evaluation studies identify conflicting results, such as: inconsistency in the best aesthetic mappings for uncertainty (Kinkeldey et al. 2014, MacEachren et al. 2005); inconsistency across different contexts (Kinkeldey et al. 2014, Hullman 2016); and inconclusive evidence for whether uncertainty is helpful at all (MacEachren et al. 2005).

There are as many explanations for these contradictions as there are contradictions themselves. Kim et al. (2019) suggests there is some interference in results from participants’ prior beliefs; Hullman (2016) notes that the irreconcilable evaluation studies are likely a product of visual heuristics, subjective probabilities, unknown participant utility functions, or a misunderstanding of statistical concepts (such as confidence intervals). Kinkeldey et al. (2014) suggest the perception of visualisation changes by audience, so we cannot expect the same results between different subpopulations. Brennen & Tuerk (2018) attributes evaluation difficulties to cognitive load from complicated uncertainty visualisations, as well as the participants’ prior experience in the topic. While these issues will certainly have some impact on our ability to synthesise information, none of them is unique to uncertainty visualisation. Rather, the issue is likely due to a disconnect between the evaluation methods used and the stated goals of each experiment, a common issue in visualisation evaluations (Vanderplas et al. 2020).

4.1. Current evaluation methods

4.1.1. Value extraction. Uncertainty visualisations are most commonly evaluated based on how accurately viewers can extract an estimate and its variance (Kinkeldey et al. 2014, Hullman et al. 2019). This is not unusual, as direct observation is the simplest way to verify that information can be accurately read from a graph (Vanderplas et al. 2020). Unfortunately, this approach doesn’t work for uncertainty visualisation. Once we ask a specific question about a statistic, that statistic becomes inferential, even if the plot was not designed for this purpose. By shifting the focus from $\hat{X}$ to $Var(\hat{X})$ or $P(\hat{X} < x)$, we end up evaluating visualisations on their ability to convey uncertainty statistics, rather than their ability to perform signal modulation. Some of the findings are obvious: participants were more accurate when reading a probability expressed as text than when they had to extract it from a graphic (Cheong et al. 2016, Savelli & Joslyn 2013). Other studies replicate existing research, such as the finding that a probability mapped to a position is more accurate than one mapped to an area (Ibrekk & Morgan 1987, Gschwandtner et al. 2016), established as part of the hierarchy of perceptual tasks more than 40 years ago (Cleveland & McGill 1984; replicated by Heer & Bostock 2010). This extends beyond simple accuracy evaluations: Sanyal et al. (2009) found that colour was more effective than size when searching for extrema in variances; we

have known that pre-attentive (subconsciously processed) aesthetics, such as colour, are more efficient for search tasks since the 1980s (Vanderplas et al. 2020). By classifying these studies as evaluations of “uncertainty visualisation” while evaluating uncertainty as a signal, we are encouraged to see successful examples of signal suppression as failure. This approach leads authors to advise against particular aesthetic mappings for uncertainty, because they cause participants to have more difficulty extracting values (Blenkinsop et al. 2000). This conclusion is antithetical to the goals of signal validation and occurs because these methods evaluate uncertainty as a signal, not as noise.

4.1.2. Trust, confidence, and risk aversion. If we cannot directly measure the effects of visual uncertainty for fear that it turns into a signal, we might instead try to measure the emotional or behavioural effect this uncertainty should have on the viewer. This seems to be the approach of many visualisation authors, as secondary benefits such as trust, confidence, and risk aversion are all frequently used in uncertainty evaluation studies (Hullman et al. 2019). Unfortunately, measuring these secondary effects often leads to confusing conclusions that simultaneously argue for and against the inclusion of uncertainty.

This is most commonly noticed in the use of trust as a measure, as several authors have commented that measuring trust, and not transparency, can lead to a questionable subtext that argues against transparency (Spiegelhalter 2017, O’Neill 2018). We see this directly play out in the visualisation literature, where surveyed visualisation authors explicitly said they didn’t include uncertainty due to the fact that they might decrease trust in their conclusions (Hullman 2020). This sentiment is also true for confidence, as Blenkinsop et al. (2000) commented that visually integrable depictions of uncertainty should be avoided, as they decrease the viewer’s confidence in their extracted data values.

Another secondary effect that is similar to trust and confidence is risk aversion. Risk aversion is an economic term used to describe an agent who would choose a random variable with a lower expected payout because it also has a lower variance. There is some disagreement on the role of risk aversion in the literature, as Hullman et al. (2019) argue uncertainty should elicit more risk aversion, whereas Wu et al. (2023) argues for less risk aversion (by proxy of suggesting rational agents as a benchmark).

Ultimately, measuring these secondary benefits has similar issues to value extraction studies, except they are slightly more confusing in their goals, leading them to simultaneously argue for and against the inclusion of uncertainty in a visualisation.

4.1.3. Alternative approaches. Often, authors understand that the effects of uncertainty are more complicated than simple value extraction. These studies indicate that accurately capturing uncertainty will be more complicated than simply avoiding value extraction or trust as a measure.

One approach is to ask indeterminate questions, such as asking participants for the “best estimate” (Ibrekk & Morgan 1987), or to select which distribution is the “furthest to the right” in a lineup (Hofmann et al. 2012). In both cases, the ground truth is based on the mean of the distribution, which is not as indeterminate as the question. This approach can lead to inconclusive results, as we are left unclear whether it was the phrasing of the question or the plot design that caused the participants to answer incorrectly.

On the other hand, questions that are incredibly specific about the distribution information can confuse the participants and induce noisy results. For example, [Hullman et al. \(2015\)](#) asked participants to compare two normally distributed groups, A and B, and had many participants say that group A was more likely to be bigger, despite group B having a higher mean. [Gschwandtner et al. \(2016\)](#) asked participants the “probability that the interval has already ended at the marked point in time?” and participants replied with the probability that the interval had already started.

4.1.4. Effective methods. This is not to say all evaluation studies fail to properly evaluate uncertainty as noise. There are several studies that ask participants to identify a particular signal that the noise is trying to obfuscate ([Kale et al. 2018](#), [Correll & Gleicher 2014](#)), which seems to be an effective method. The only problem with these studies is that most uncertainty visualisation methods are not integrated into *the grammar of graphics* ([Wilkinson 2005](#), [Wickham 2010](#)), so we regularly see comparisons between disparate plots that would never be considered substitutes for one another outside the artificial “uncertainty visualisation” framework they are placed within. For example, [Kale et al. \(2018\)](#) compared static bar charts with error bars to a bar chart with animated samples, meaning that any difference in participants’ ability to read the plot could be due to the statistic (confidence interval versus sample), the geometry (bars versus intervals), or the use of animation (static versus animated plots). This issue was rectified in their second experiment, where they compared overlaid and animated samples, giving us an insight into the types of visualisations that are appropriate to compare in uncertainty visualisation experiments. This means that even when evaluations *are* done correctly, there is no generalisable theory we can take from the results.

The point here is not to accuse the authors of poor academic rigour. The papers are (usually) logically consistent and well-formulated pieces of work. Rather, the point is to illustrate that evaluating uncertainty as noise is surprisingly difficult. Designing tests for signal validation will require a formalisation of uncertainty within the grammar of graphics, as well as improved evaluation methods.

4.2. Implicit Hypothesis Testing

The main problem with current uncertainty visualisation evaluations is that they often require explicit (or convoluted) questions about the variance. Asking direct questions about the statistics or outcomes is not an explicit requirement of visualisation evaluations. In their review of testing statistical graphics, [Vanderplas et al. \(2020\)](#) drew a distinction between explicit tests, where participants are asked direct questions about specific features of a plot, and implicit testing, where users identify both the purpose and function of the plot. The lineup protocol is the most salient example of the implicit approach. Lineups are a confirmatory visualisation tool where participants are shown a set of M plots, and asked to identify the plot that is the “most different”, leaving participants to decide what “most different” means to them, even if it is not what the authors intended ([VanderPlas & Hofmann 2017](#)). The implicit test does not limit the versatility of the approach, with the lineup being used to evaluate the effectiveness of different types of plots ([Hofmann et al. 2012](#)), colour palettes ([Reda & Szafr 2021](#)), and design decisions ([VanderPlas & Hofmann 2017](#)).

Lineup protocols are not only useful for implicit testing: they also have parallels to hypothesis

testing that can be leveraged in uncertainty visualisation. The concept of signal validation is, at its core, an assertion of statistical validity: the visibility of signals should be directly proportional to p -values or some equivalent measure. This comparison is not new in uncertainty visualisation, where parallels have been drawn to frequentist statistics by [Correll & Gleicher \(2014\)](#), who compared results to Cohen's D , and to Bayesian statistics by [Kim et al. \(2019\)](#), who evaluated plots based on their impact on the users' prior beliefs. The comparison to hypothesis testing is far more natural for the lineup protocol, which has a visual test statistic and can be compared to results from Bayesian analysis ([VanderPlas & Hofmann 2017](#)) or standard statistical tests using power curves ([Majumder et al. 2013](#), [Li et al. 2024](#)). The connections between lineups and uncertainty visualisation are numerous and have been previously identified in the development of Hypothetical Outcome Plots ([Hullman et al. 2015](#)).

The lineup protocol and uncertainty visualisations are similar: lineups were designed for checking if perceived patterns are real or merely the result of chance ([Buja et al. 2009](#), [Wickham et al. 2010](#)). As both approaches are attempting to do the same thing, it is likely that we are unable to leverage the lineup protocol directly to evaluate uncertainty visualisation, but a new evaluation methodology should try to learn from the success of the lineup approach. Designing an implicit testing method for uncertainty visualisation that allows us to draw parallels to standard notions of statistical significance would solve many of the issues with the current evaluation approaches.

5. Conclusions and Future Work

This paper examines the literature and provides suggestions for a structural framework to support uncertainty visualisation. Particularly, we propose that uncertainty visualisation should act as a signal modulation, suppressing weak signals, and amplifying strong signals. Using this framework, there are four key gaps identified by this review.

Experimental practices on uncertainty visualisation need standards. Some existing evaluation experiments treat uncertainty as a signal, while others treat uncertainty as noise. As a result, it is difficult to combine results from papers to get a meaningful sense of how uncertainty information is understood by a viewer. Researchers need to ensure that when they identify the motivation behind their visualisation technique, their evaluation methods align with the stated goals of the paper.

Experimental methods that evaluate uncertainty as noise need to be developed. Research into separability and integrability of signal and noise is of particular interest to uncertainty visualisation, as it allows assessment of the interference between the two. When designing experiments, authors often choose aesthetics that are visually distinguishable; uncertainty visualisation authors should consider doing the opposite.

Uncertainty needs to be formalised within the grammar of graphics. Some of this formalisation was done by [Kay \(2023\)](#), but it focuses only on the visualisation of univariate distributions. Giving authors the ability to describe uncertainty visualisations in terms of statistics, geometries, and aesthetics will support evaluation experiments that can build towards a cohesive theory of visualising uncertainty.

Software that allows users to easily perform signal validation is needed. Existing uncertainty

visualisation methods view a distribution as its own object, and there are no software options treating “an uncertainty visualisation as a function of an existing visualisation” philosophy.

signal modulation is an undeveloped area of visualisation research, and developing methods for the practice may require us to challenge our entire notion of what makes a good visualisation.

Reproducibility

The R packages were used for this work were: `tidyverse` (Wickham et al. 2019), `RColorBrewer` (Neuwirth 2022), `scales` (Wickham et al. 2023), `sf` (Pebesma & Bivand 2023), `urbanmapr` (Strochak et al. 2024), `flextable` (Gohel & Skintzos 2024), `colorspace` (Stauffer et al. 2009), `ggdist` (Kay 2023), `ggdibbler` (Mason et al. 2026), `patchwork` (Pedersen 2025), `distributional` (O’Hara-Wild et al. 2024), `ggthemes` (Arnold 2024), `broom` (Robinson et al. 2026), and `rgeos` (Bivand & Rundel 2023). The GitHub repository for this paper can be found at <https://github.com/harriet-mason/ARSA-UncertaintyLitReview>, which contains the files required to reproduce this article in full.

DISCLOSURE STATEMENT

The authors are not aware of any affiliations, memberships, funding, or financial holdings that might be perceived as affecting the objectivity of this review.

ACKNOWLEDGMENTS

The first author of this paper is supported in part by a scholarship from the Australian Energy Market Operator. This research was supported by the Commonwealth through an Australian Government Research Training Program Scholarship [DOI: <https://doi.org/10.82133/C42F-K220>].

LITERATURE CITED

- Anscombe FJ. 1973. Graphs in statistical analysis. *The American Statistician* 27(1):17–21
- Arnold JB. 2024. ggthemes: Extra themes, scales and geoms for 'ggplot2'. R package version 5.1.0
- Begg SH, Welsh MB, Bratvold RB. 2014. Uncertainty vs. variability: What’s the difference and why is it important?, In *SPE Hydrocarbon Economics and Evaluation Symposium*, vol. SPE Hydrocarbon Economics and Evaluation Symposium
- Belsley D, Kuh E, Welsch R. 1980. Regression diagnostics: Identifying influential data and sources of collinearity. Wiley Series in Probability and Statistics. New York: Wiley
- Benjamin DM, Budescu DV. 2018. The role of type and source of uncertainty on the processing of climate models projections. *Frontiers in Psychology* 9(MAR):1–17
- Bivand R, Rundel C. 2023. rgeos: Interface to geometry engine - open source ('geos'). R package version 0.6-3
- Blenkinsop S, Fisher P, Bastin L, Wood J. 2000. Evaluating the perception of uncertainty in alternative visualization strategies. *Cartographica* 37(1):1–13
- Boukhelifa N, Bezerianos A, Isenberg T, Fekete JD. 2012. Evaluating sketchiness as a visual variable for the depiction of qualitative uncertainty. *IEEE Transactions on Visualization and Computer Graphics* 18(12):2769–2778

- Boukhelifa N, Perrin ME, Huron S, Eagan J. 2017. How data workers cope with uncertainty: A task characterisation study, In *Proceedings of the 2017 CHI Conference on Human Factors in Computing Systems*, CHI '17, p. 3645–3656, New York, NY, USA: Association for Computing Machinery
- Brennen A, Tuerk S. 2018. An instrument for evaluating uncertainty visualization techniques. *Conference on Human Factors in Computing Systems - Proceedings* 2018-April:1–6 ISBN: 9781450356206
- Buja A, Cook D, Hofmann H, Lawrence M, Lee EK, et al. 2009. Statistical inference for exploratory data analysis and model diagnostics. *Philosophical Transactions of the Royal Society A: Mathematical, Physical and Engineering Sciences* 367(1906):4361–4383
- Chakraborty S, Kiefer P, Raubal M. 2024. The influence of uncertainty visualization on cognitive load in a safety- and time-critical decision-making task. *International Journal of Geographical Information Science* 38(8):1583–1610
- Chatfield C. 1985. The initial examination of data. *Journal of the Royal Statistical Society. Series A (General)* 148(3):214–253
- Cheong L, Bleisch S, Kealy A, Tolhurst K, Wilkening T, Duckham M. 2016. Evaluating the impact of visualization of wildfire hazard upon decision-making under uncertainty. *International Journal of Geographical Information Science* 30(7):1377–1404
- Cleveland WS, McGill R. 1984. Graphical perception: Theory, experimentation, and application to the development of graphical methods. *Journal of the American Statistical Association* 79(387):531–554
- Correll M, Gleicher M. 2014. Error bars considered harmful: Exploring alternate encodings for mean and error. *IEEE Transactions on Visualization and Computer Graphics* 20(12):2142–2151
- Correll M, Moritz D, Heer J. 2018. Value-suppressing uncertainty palettes. *Conference on Human Factors in Computing Systems - Proceedings* 2018-April:1–11
- Fischhoff B, Davis AL. 2014. Communicating scientific uncertainty. *Proceedings of the National Academy of Sciences of the United States of America* 111:13664–13671
- Franconeri SL. 2021. Three perceptual tools for seeing and understanding visualized data. *Current Directions in Psychological Science* 30(5):367–375
- Gohel D, Skintzos P. 2024. flextable: Functions for tabular reporting. R package version 0.9.6
- Goldstein DG, Rothschild D. 2014. Lay understanding of probability distributions. *Judgment and Decision Making* 9(1):1–14
- Griethe H, Schumann H. 2006. The visualization of uncertain data: Methods and problems, In *SimVis*, vol. 6, pp. 143–156
- Gschwandtner T, Bögl M, Federico P, Miksch S. 2016. Visual encodings of temporal uncertainty: A comparative user study. *IEEE Transactions on Visualization and Computer Graphics* 22(1):539–548
- Guo Z, Kale A, Kay M, Hullman J. 2025. VMC: A grammar for visualizing statistical model checks. *IEEE Transactions on Visualization and Computer Graphics* 31(1):798–808
- Gustafson A, Rice RE. 2019. The effects of uncertainty frames in three science communication topics. *Science Communication* 41(6):679–706
- Hadjimichael A, Schlumberger J, Haasnoot M. 2024. Data visualisation for decision making under deep uncertainty: current challenges and opportunities. *Environmental Research Letters* 19(11):111011
- Heer J, Bostock M. 2010. Crowdsourcing graphical perception: using mechanical turk to assess visualization design, In *Proceedings of the SIGCHI Conference on Human Factors in Computing Systems*, CHI '10, p. 203–212, New York, NY, USA: Association for Computing Machinery
- Henderson, Velleman. 1981. Building multiple regression models interactively. *Biometrics* 37:391–411
- Hofmann H, Follett L, Majumder M, Cook D. 2012. Graphical tests for power comparison of competing designs. *IEEE Transactions on Visualization and Computer Graphics* 18(12):2441–2448
- Huebner M, Vach W, le Cessie S. 2016. A systematic approach to initial data analysis is good research practice. *The Journal of Thoracic and Cardiovascular Surgery* 151(1):25–27
- Hullman J. 2016. Why evaluating uncertainty visualization is error prone. *ACM International*

- Conference Proceeding Series* 24-October:143–151
- Hullman J. 2020. Why authors don't visualize uncertainty. *IEEE Transactions on Visualization and Computer Graphics* 26(1):130–139
- Hullman J, Gelman A. 2021. Designing for interactive exploratory data analysis requires theories of graphical inference. *Harvard Data Science Review* :1–70
- Hullman J, Kay M, Kim YS, Shrestha S. 2018. Imagining replications: Graphical prediction discrete visualizations improve recall estimation of effect uncertainty. *IEEE Transactions on Visualization and Computer Graphics* 24(1):446–456
- Hullman J, Qiao X, Correll M, Kale A, Kay M. 2019. In pursuit of error: A survey of uncertainty visualization evaluation. *IEEE Transactions on Visualization and Computer Graphics* 25(1):903–913
- Hullman J, Resnick P, Adar E. 2015. Hypothetical outcome plots outperform error bars and violin plots for inferences about reliability of variable ordering. *PLoS ONE* 10(11)
- Hyndman R, Athanasopoulos G. 2021. Forecasting: principles and practice, 3rd edition. Melbourne, Australia: OTexts
- Ibrekk H, Morgan MG. 1987. Graphical communication of uncertain quantities to nontechnical people. *Risk Analysis* 7(4):519–529
- Kale A, Kay M, Hullman J. 2021. Visual reasoning strategies for effect size judgments and decisions. *IEEE Transactions on Visualization and Computer Graphics* 27(2):272–282
- Kale A, Nguyen F, Kay M, Hullman J. 2018. Hypothetical outcome plots help untrained observers judge trends in ambiguous data. *IEEE transactions on visualization and computer graphics* 25(1):892–902
- Kay M. 2019. How much value should an uncertainty palette suppress if an uncertainty palette should suppress value? Statistical and perceptual perspectives
- Kay M. 2023. ggdist: Visualizations of distributions and uncertainty in the grammar of graphics. *IEEE Transactions on Visualization and Computer Graphics* 30(1):414–424
- Kim YS, Walls LA, Krafft P, Hullman J. 2019. A bayesian cognition approach to improve data visualization. *Conference on Human Factors in Computing Systems - Proceedings* :1–14
- Kinkeldey C, MacEachren AM, Schiewe J. 2014. How to assess visual communication of uncertainty? A systematic review of geospatial uncertainty visualisation user studies. *Cartographic Journal* 51(4):372–386
- Koonchanok R, Tawde GY, Narayanasamy GR, Walimbe S, Reda K. 2023. Visual belief elicitation reduces the incidence of false discovery. *Conference on Human Factors in Computing Systems - Proceedings*
- Kuhnert PM, Pagendam DE, Bartley R, Gladish DW, Lewis SE, Bainbridge ZT. 2018. Making management decisions in the face of uncertainty: A case study using the Burdekin catchment in the Great Barrier Reef. *Marine and Freshwater Research* 69(8):1187–1200
- Lee C, Yang T, Inchoco GD, Jones GM, Satyanarayan A. 2021. Viral visualizations: How coronavirus skeptics use orthodox data practices to promote unorthodox science online, In *Proceedings of the 2021 CHI conference on human factors in computing systems*, pp. 1–18
- Li W, Cook D, Tanaka E, VanderPlas S. 2024. A plot is worth a thousand tests: Assessing residual diagnostics with the lineup protocol. *Journal of Computational and Graphical Statistics* :1–19
- Locke S, D'Agostino McGowan L. 2018. datasauRus: Datasets from the datasaurus dozen. R package version 0.1.4
- Lucchesi L, Kuhnert P, Wikle C. 2021. Vizumap: an R package for visualising uncertainty in spatial data. *Journal of Open Source Software* 6(59):2409
- MacEachren AM. 1992. Visualizing uncertain information. *Cartographic Perspectives* (13):10–19
- MacEachren AM, Robinson A, Hopper S, Gardner S, Murray R, et al. 2005. Visualizing geospatial information uncertainty: What we know and what we need to know. *Cartography and Geographic Information Science* 32(3):139–160ISBN: 1523040054738
- MacEachren AM, Roth RE, O'Brien J, Li B, Swingle D, Gahegan M. 2012. Visual semiotics &

- uncertainty visualization: An empirical study. *IEEE Transactions on Visualization and Computer Graphics* 18(12):2496–2505
- Majumder M, Hofmann H, Cook D. 2013. Validation of visual statistical inference, applied to linear models. *Journal of the American Statistical Association* 108(503):942–956
- Manski CF. 2020. The lure of incredible certitude. *Economics and Philosophy* 36(2):216–245
- Mason H, Cook D, Goodwin S, VanderPlas S. 2026. ggdbbler: Add uncertainty to data visualisations
- Matejka J, Fitzmaurice G. 2017. Same stats, different graphs: Generating datasets with varied appearance and identical statistics through simulated annealing, In *Proceedings of the 2017 CHI Conference on Human Factors in Computing Systems*, CHI '17, p. 1290–1294, New York, NY, USA: Association for Computing Machinery
- Meng XL. 2014. A trio of inference problems that could win you a nobel prize in statistics (if you help fund it). *Past, Present, and Future of Statistical Science* :537–562
- Moritz D, Fisher D, Ding B, Wang C. 2017. Trust, but verify: Optimistic visualizations of approximate queries for exploring big data, In *Proceedings of the 2017 CHI conference on human factors in computing systems*, pp. 2904–2915
- Ndlovu A, Shrestha H, Harrison LT. 2023. Taken by surprise? evaluating how bayesian surprise & suppression influences peoples' takeaways in map visualizations, In *2023 IEEE Visualization and Visual Analytics (VIS)*, pp. 136–140, IEEE
- Neuwirth E. 2022. RColorBrewer: ColorBrewer palettes. R package version 1.1-3
- O'Hara-Wild M, Kay M, Hayes A, Hyndman R. 2024. distributional: Vectorised probability distributions. R package version 0.5.0
- Olston C, Mackinlay JD. 2002. Visualizing data with bounded uncertainty. *Proceedings - IEEE Symposium on Information Visualization, INFO VIS 2002-Janua*:37–40
- O'Neill O. 2018. Linking trust to trustworthiness. *International Journal of Philosophical Studies* 26(2):293–300
- Otsuka J. 2023. Thinking about statistics: The philosophical foundations. New York: Routledge, 1st ed.
- Padilla L, Hosseinpour H, Fygenon R, Howell J, Chunara R, Bertini E. 2022a. Impact of COVID-19 forecast visualizations on pandemic risk perceptions. *Scientific Reports 2022 12:1* 12(1):1–14
- Padilla L, Kay M, Hullman J. 2022b. Uncertainty visualization. In *Computational Statistics in Data Science*, eds. WW Piegorsch, RA Levine, HH Zhang, TCM Lee, chap. 22. Hoboken, NJ: John Wiley & Sons, 405–426
- Padilla L, Powell M, Kay M, Hullman J. 2021. Uncertain about uncertainty: How qualitative expressions of forecaster confidence impact decision-making with uncertainty visualizations. *Frontiers in Psychology* 11
- Pang AT, Wittenbrink CM, Lodha SK. 1997. Approaches to uncertainty visualization. *Visual Computer* 13(8):370–390
- Pebesma E, Bivand R. 2023. Spatial data science: With applications in R. Chapman and Hall/CRC
- Pedersen TL. 2025. patchwork: The composer of plots. R package version 1.3.2
- Peña-Araya V, Fontaine CM, Wei X, Delpech G, Bezerianos A. 2025. Uncertainty in science is malleable. advocating for user-agency in defining uncertainty in visualizations: a case study in geology, In *Proceedings of the 2025 CHI Conference on Human Factors in Computing Systems*, pp. 1–18
- Pham B, Streit A, Brown R. 2009. Visualization of information uncertainty: Progress and challenges. In *Advanced Information and Knowledge Processing*, vol. 36. Springer-Verlag London Ltd, 19–48
- Potter K, Kniss J, Riesenfeld R, Johnson CR. 2010. Visualizing summary statistics and uncertainty. *Computer Graphics Forum* 29(3):823–832
- Reda K, Szafr DA. 2021. Rainbows revisited: Modeling effective colormap design for graphical inference. *IEEE Transactions on Visualization and Computer Graphics* 27(2):1032–1042
- Robinson D, Hayes A, Couch S, Hvitfeldt E. 2026. broom: Convert statistical objects into tidy tibbles. R package version 1.0.12

- Sanyal J, Zhang S, Bhattacharya G, Amburn P, Moorhead R.J. 2009. A user study to compare four uncertainty visualization methods for 1D and 2D datasets. *IEEE Transactions on Visualization and Computer Graphics* 15(6):1209–1218
- Sarma A, Guo S, Hoffswell J, Rossi R, Du F, et al. 2023. Evaluating the use of uncertainty visualisations for imputations of data missing at random in scatterplots. *IEEE Transactions on Visualization and Computer Graphics* 29(1):602–612
- Sarma A, Pu X, Cui Y, Correll M, Brown ET, Kay M. 2024. Odds and insights: Decision quality in exploratory data analysis under uncertainty, In *Proceedings of the CHI Conference on Human Factors in Computing Systems*, CHI '24. New York, NY, USA: Association for Computing Machinery
- Savelli S, Joslyn S. 2013. The advantages of predictive interval forecasts for non-expert users and the impact of visualizations. *Applied Cognitive Psychology* 27(4):527–541
- Savvides R, Henelius A, Oikarinen E, Puolamäki K. 2019. Significance of patterns in data visualisations, In *Proceedings of the 25th ACM SIGKDD International Conference on Knowledge Discovery & Data Mining*, pp. 1509–1517, Anchorage AK USA: Association for Computing Machinery
- Smart S, Szafr DA. 2019. Measuring the separability of shape, size, and color in scatterplots. *Conference on Human Factors in Computing Systems - Proceedings* :1–14
- Spiegelhalter D. 2017. Risk and uncertainty communication. *Annual Review of Statistics and Its Application* 4:31–60
- Stauffer R, Mayr GJ, Dabernig M, Zeileis A. 2009. Somewhere over the rainbow: How to make effective use of colors in meteorological visualizations. *Bulletin of the American Meteorological Society* 96(2):203–216
- Strochak S, Ueyama K, Williams A. 2024. urbnmapr: State and county shapefiles in sf and tibble format. R package version 0.0.0.9002
- Thomson J, Hetzler E, MacEachren A, Gahegan M, Pavel M. 2005. A typology for visualizing uncertainty. *Visualization and Data Analysis 2005* 5669(March 2005):146
- Tierney N, Cook D. 2023. Expanding tidy data principles to facilitate missing data exploration, visualization and assessment of imputations. *Journal of Statistical Software* 105(7):1–31
- Tukey JW, et al. 1977. Exploratory data analysis, vol. 2. Springer
- Vanderplas S, Cook D, Hofmann H. 2020. Testing statistical charts: What makes a good graph? *Annual Review of Statistics and Its Application* 7(1):61–88
- VanderPlas S, Hofmann H. 2015. Signs of the sine illusion—why we need to care. *Journal of Computational and Graphical Statistics* 24(4):1170–1190
- VanderPlas S, Hofmann H. 2017. Clusters beat trend!? Testing feature hierarchy in statistical graphics. *Journal of Computational and Graphical Statistics* 26(2):231–242
- Walker WE, Harremoes P, Rotmans J, Van Der Sluijs JP, Van Asselt MBA, et al. 2003. Defining uncertainty. *Integrated Assessment* 4(1):5–17
- Wallsten TS, Budescu DV, Erev I, Diederich A. 1997. Evaluating and combining subjective probability estimates. *Journal of Behavioral Decision Making* 10(3):243–268
- Wickham H. 2010. A layered grammar of graphics. *Journal of Computational and Graphical Statistics* 19(1):3–28
- Wickham H, Averick M, Bryan J, Chang W, D’Agostino McGowan L, et al. 2019. Welcome to the tidyverse. *Journal of Open Source Software* 4(43):1686
- Wickham H, Cook D, Hofmann H, Buja A. 2010. Graphical inference for infovis. *IEEE Transactions on Visualization and Computer Graphics* 16:973–979
- Wickham H, Hofmann H. 2011. Product plots. *IEEE Transactions on Visualization and Computer Graphics* 17(12):2223–2230
- Wickham H, Pedersen TL, Seidel D. 2023. scales: Scale functions for visualization. R package version 1.3.0
- Wilkinson L. 2005. The grammar of graphics. Berlin, Heidelberg: Springer-Verlag

- Wu Y, Guo Z, Mamakos M, Hartline J, Hullman J. 2023. The rational agent benchmark for data visualization
- Yang F, Cai M, Mortenson C, Fakhari H, Lokmanoglu AD, et al. 2023. Swaying the public? Impacts of election forecast visualizations on emotion, trust, and intention in the 2022 us midterms. *IEEE Transactions on Visualization and Computer Graphics* 30(1):23–33
- Zhao J, Wang Y, Mancenido MV, Chiou EK, Maciejewski R. 2023. Evaluating the impact of uncertainty visualization on model reliance. *IEEE Transactions on Visualization and Computer Graphics* 30(7):4093–4107